

BEOMYEOL JEON

bj2@illinois.edu ◇ beomyeol.github.io ◇ in beomyeol ◇ beomyeol

EDUCATION

- University of Illinois Urbana-Champaign**, Urbana, IL May 2024 (Expected)
Ph.D in Computer Science. GPA: 3.97/4.0
- Seoul National University**, Seoul, Korea Feb 2016
B.S. in Computer Science and Engineering (*summa cum laude*). GPA: 4.0/4.3
- The University of Texas at Austin**, Austin, TX Jan 2015 – May 2015
Undergraduate Exchange Student in Computer Science. GPA: 3.91/4.0

WORK EXPERIENCE

- University of Illinois Urbana-Champaign** Aug 2016 – Present
Graduate Research Assistant. Advisor: Prof. Indranil Gupta Urbana, IL
- Designed and built a fast device placement (model parallelism) system for ML graphs (e.g., Transformers, LLMs, etc.) via an algorithmic approach. Our system is integrated with TensorFlow and generates expert-level quality placements **654×–206K× faster** than the state-of-the-art ML-based approaches (published in SoCC '20) [code].
 - Designed and built an ML inference framework that autoscales inference jobs with time-varying workloads via probabilistic time-series prediction and relaxed optimization (in collaboration with IBM Research). Our system is integrated with Ray Serve on a Kubernetes cluster and achieves **1.7×–22× lower SLO violations** than state-of-the-art systems.
 - Designing and building an autotuning Graph Neural Network (GNN) training framework on AWS Lambda via simulation and various optimization techniques (e.g., gray-box optimization, RL, etc.)
- Nokia Bell Labs** Jun 2019 – Aug 2019
Research Intern. Mentor: Dr. Muntasir Raihan Rahman and Dr. Anwar Walid Murray Hill, NJ
- Designed a privacy-preserving decentralized aggregation protocol for federated learning by controlling communication patterns in ADMM-based aggregation (published in FOGML '21)[code].
- Google** May 2018 – Aug 2018
Software Engineering Intern. BigQuery Team. Mentor: Dr. Mohsen Vakilian Kirkland, WA
- Built an autotuner for BigQuery system configuration parameters to improve performance via black-box optimization.
- Google** May 2017 – Aug 2017
Software Engineering Intern. Earth Engine Team. Mentor: Dr. Hector Gonzalez Mountain View, CA
- Designed and built a prototype framework for large-scale GIS batch computation on top of Google Cloud Dataflow.
- Somansa** Jan 2011 – Dec 2013
Researcher [Full-time, in fulfillment of mandatory military service duty] Seoul, Korea
- Developed a network data loss prevention (DLP) solution, Mail-i, and a database audit and protection (DAP) solution, DB-i, by analyzing network packets through reverse engineering.

PUBLICATIONS

- Beomyeol Jeon**, Yongjoo Park, Indranil Gupta. *Automating Resource Allocation for Graph Neural Network Training on Serverless Frameworks*, Currently Under Preparation.
- Beomyeol Jeon**, Chen Wang, Diana Arroyo, Alaa Youssef, Indranil Gupta. *SLO-aware ML Inference Autoscaler for Fixed-Size On-Premises Clusters*, Under Review at a Conference.
- Beomyeol Jeon***, S M Ferdous*, Muntasir Raihan Rahman, Anwar Walid. *Privacy-preserving Decentralized Aggregation for Federated Learning*. FOGML 2021. [link, code, extended version]

Beomyeol Jeon, Linda Cai, Pallavi Srivastava, Jintao Jiang, Xiaolan Ke, Yitao Meng, Cong Xie, Indranil Gupta. *Bacchi: Fast Device Placement of Machine Learning Graphs*. SoCC 2020. [link, code, extended version]

Woo-Yeon Lee, ... **Beomyeol Jeon**, et al. *Automating System Configuration of Distributed Machine Learning*. ICDCS 2019. [link, code]

Byung-Gon Chun, ..., **Beomyeol Jeon**, et al. *Apache REEF: Retainable Evaluator Execution Framework*. ACM TOCS, Volume 35 Issue 2. [link]

Byung-Gon Chun, Brian Cho, **Beomyeol Jeon**, Joo Seong Jeong, Gunhee Kim, Joo Yeon Kim, Woo-Yeon Lee, Yun Seong Lee, Markus Weimer, Gyeong-In Yu. *Dolphin: Runtime Optimization for Distributed Machine Learning*. ICML MLSys 2016 workshop. [link]

SELECTED TECHNICAL SKILLS

Languages	C/C++, Python, Go, Java, C#
Frameworks & Platforms	PyTorch, TensorFlow, Ray, Kubernetes, Docker, AWS EC2, AWS Lambda
Libraries	Boost, Eigen, OpenMP, CUDA, ZeroMQ, Protocol Buffers, Scikit-learn, TBB
Databases	Redis, Microsoft SQL Server
Tools	Git, Vim, Visual Studio Code, Bash/Zsh, GitHub, Pandas